

DS Lecture 7.2

Leader Election

October 24, 2022

Michele Albano
mialb@cs.aau.dk

DEIS
Aalborg University
Denmark



AALBORG UNIVERSITY
DENMARK

Based on slides by Peter G. Jensen, AAU.

What is it?



What is leader election?

Leader election algorithms ensure that one, and only one process is elected as the leader in the event of lack of a leader



What is it?

Examples

- ▶ Algorithms with coordinator
 - ▶ Mutex?
- ▶ Distributed replication
- ▶ DNS-servers in case of network partition
- ▶ Failover mechanism for crashes
- ▶ Parliaments?

What is it?



Big Question:

We detected a crash of the leader, now what?

Agenda



Requirements

Assumptions

Chang-Roberts

Bully Algorithm

Requirements

Leader Election

Let $P = \{p_1, \dots, p_n\}$ be a set of processes and let $L(p_i) \in P \cup \{\perp\}$ be the leader as seen from a process p_i

1. Safety

- ▶ either $L(p_i) = \perp$ or $L(p_i) = p_j$
 - ▶ p_j is the non-crashed process with largest identifier

2. Liveness

- ▶ All process participate, and eventually
 - ▶ p_i crashes, or
 - ▶ $L(p_i) \neq \perp$

Note

Processes can crash during the election

Assumptions



- ▶ Processes stay dead
- ▶ Crashes are reliably detected
- ▶ Identifiers are unique

Chang Roberts

Idea

- ▶ Pass token of largest ID in a ring
- ▶ Basic algorithm in election
 - ▶ Forward ID to “next” if higher than own,
 - ▶ Forward own ID to “next” otherwise.
 - ▶ Only one active election
- ▶ Two message types:
 - ▶ *election* to vote. If p_i receives it and is not "participating", it starts participating
 - ▶ *elected* to declare who won. If p_i is participating and receives its own ID, it means it won, becomes "non-participating" and send the *elected* message

Chang-Roberts - Code

```

1  def run(self):
2      while True:
3          nxt = (self.index() + 1) % self.
            number_of_devices()
4          if not self._participated:
5              self.medium().send(
6                  Vote(self.index(), nxt, self.index(),
                        False))
7              self._participated = True
8              ingoing = self.medium().receive()
9              if ingoing is not None:
10                 if ingoing.vote() == self.index():
11                     if not ingoing.decided():
12                         self.medium().send(
13                             Vote(self.index(), nxt, self.
                                    index(), True))
14             else:
15                 self._leader = self.index()
16                 return # this device is the new
                        leader

```

```

17         elif ingoing.vote() < self.index():
18             continue
19         elif ingoing.vote() > self.index():
20             # forward the message
21             self.medium().send(
22                 Vote(self.index(), nxt, ingoing.vote(),
                        ingoing.decided()))
23         if ingoing.decided():
24             self._leader = ingoing.vote()
25             return
26         self.medium().wait_for_next_round()

```

Properties

Chang Roberts



- ▶ Safe
- ▶ Live
- ▶ $3N - 1$ messages per election

Crashes?

Can be overcome if reliably detected!



Bully Algorithm

Idea

- ▶ Bully election requests into silence
- ▶ Priority by ID
- ▶ Basic algorithm in election
 - ▶ Send “shut up” to lower IDs
 - ▶ Higher IDs will “answer”
 - ▶ Finally, highest alive ID broadcasts itself

NOTE!

Depends on Synchronous Behavior (timeout to detect crashes)

```

class Bully (Device):
    def __init__ (...):
        super().__init__(index, number_of_devices, medium)
        self._leader = None
        self._shut_up = False
        self._election = False

    def run(self):
        first_round = True
        while True:
            got_input = False
            if not self._shut_up and not self._election:
                self.start_election()
            new_election = False
            while True:
                ingoing = self.medium().receive()
                if ingoing is not None:
                    got_input = True
                    if ingoing.vote() < self.index():
                        self.medium().send(
                            Vote(self.index(), ingoing.source, self
                                .index(), self.largest()))
                        new_election = True
                else:
                    self._shut_up = True
                    if ingoing.decided():
                        self._leader = ingoing.vote()
                        return
            else: break

```

```

if not self._shut_up and not self._election and
    new_election:
        self.start_selection()
if not got_input and not first_round:
    if self._election:
        if self._shut_up:
            self._shut_up = False
            self.start_election()
        else:
            for id in self.medium().ids():
                if id != self.index():
                    self.medium().send(Vote(self.index
(), id, self.index(), True))
                    self._leader = self.index()
                    return
            self.medium().wait_for_next_round()
            first_round = False

def largest(self):
    return self.index() == max(self.medium().ids())

def start_election(self):
    if not self._election:
        self._election = True
        for id in self.medium().ids():
            if id > self.index():
                self.medium().send(Vote(self.index(), id, self
.index(), self.largest()))

```

Properties

Bully



- ▶ Safe & Live, assuming
 - ▶ Unique IDs
 - ▶ Failure detection is reliable
- ▶ Best-case: $N - 2$ messages pr election
- ▶ Worst-case: $O(N^2)$ messages (the process with the lowest ID starts the election)
- ▶ Election-time: 2 rounds (assuming HW multicast)

Properties

Bully

Beware

Safety is broken if

- ▶ too tight deadline,
- ▶ process IDs reappear, or
- ▶ system is not synchronous.

Requires

- ▶ requires synchronous system , and
- ▶ ordering of messages